Original synthetic VAL862: SFT versus all baseline references

Supplementary: equal weight per configuration | 862 configurations / 500 DAGs | 77 full-panel evaluations; no smoothing

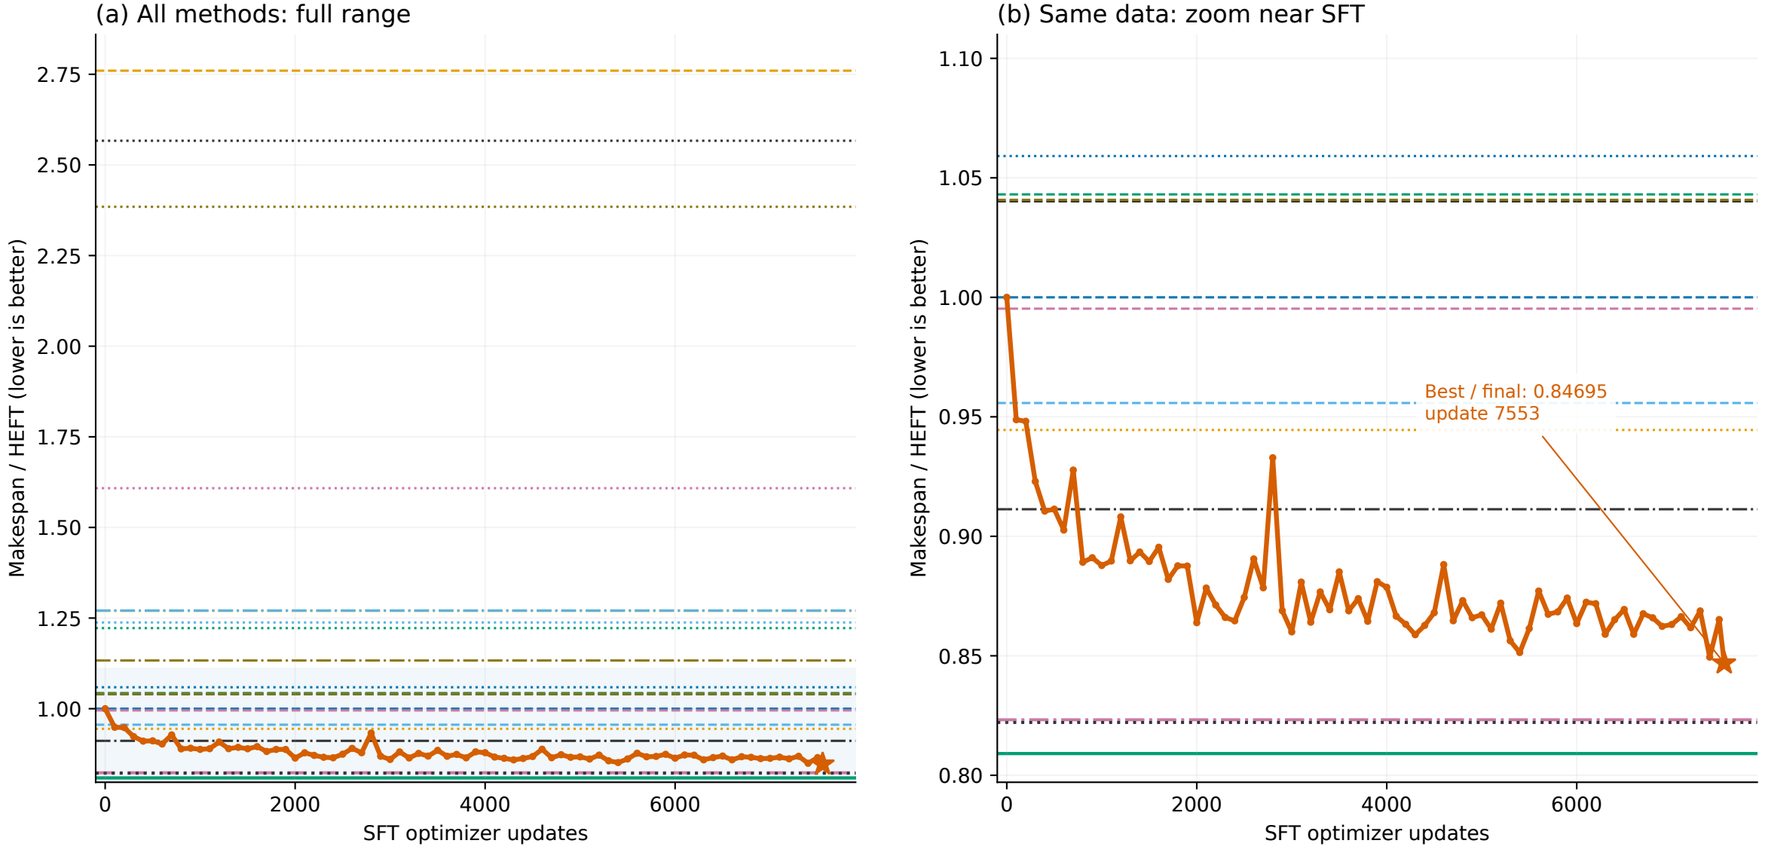

Method / line key	C / HEFT
SFT (ours; final)	**0.84695**
Teacher64+BHEFT-v2	0.80908
Teacher64 (legacy)	0.82210
Portfolio11 (no search)	0.82323
GiPH / HEFT start / 2n	0.91132
BHEFT	0.94449
MSLBL-adapted	0.95578
JoinLookahead	0.99525
HEFT	1.00000
PEFT-price0.1	1.04017
PEFT-price0.3	1.04080
PEFT	1.04307
PEFT-price1	1.05912
GiPH / random start / 2n	1.13294
PEFT-price3	1.22225
Greedy	1.23782
DAG-CAI-DR / BiGAT	1.27004
DAG-CAI-DR / Graphormer	1.27094
PEFT-price10	1.60829
BUDA	2.38449
RandomMean5	2.56667
CostFirst	2.75993

SFT: original 17,261 TRAIN configurations; Teacher64+BHEFT-v2 labels; seed 28; 7 epochs / 7,553 updates.

Evaluations: update 0, every 100 updates, and final update 7,553. Each baseline is a fixed horizontal reference.

Rules use native ordering; Portfolio11 selects fixed-HEFT mapping replays. Learned controls are local adaptations (seed 11).

CAI: 7 training epochs, historical selected epoch; GiPH: 1 training epoch / 16 moves per graph, up to 2n evaluation moves.